

# AI CompanyOps Multi-Agent System

## *Project Proposal*

### **Project Members**

Chen Zihao

Jin Yanyu

Li Lin

Yang Miao

Zheng Jiecheng

# Overview

---

## Background and Business Problem

In modern organizations, business operations require coordinated contributions from multiple specialized departments, including product management, engineering, marketing, finance, and quality assurance. Designing comprehensive project proposals or business solutions typically demands significant human effort, time, and cross-functional coordination. Teams must gather domain expertise, align on strategic direction, and synthesize diverse outputs into a coherent plan. This process is often slow, resource-intensive, and prone to inconsistencies arising from siloed communication between departments.

The AI CompanyOps Multi-Agent System addresses this challenge by simulating a fully operational company through a coordinated network of AI agents, each representing a specific organizational department. The system enables users to submit a high-level business request and receive a comprehensive, multi-perspective project proposal without requiring manual coordination across multiple subject-matter experts.

## Proposed System

The system is built around five specialized AI agents, orchestrated by a central CEO Agent. When a user submits a business task, the CEO Agent decomposes the request into domain-specific subtasks and coordinates the activities of the Product Agent, Engineering Agent, Finance Agent, and QA Agent. Each agent operates autonomously within its defined scope, producing structured outputs that are passed through the system according to a managed workflow. The QA Agent reviews all intermediate outputs for consistency and completeness before the CEO Agent consolidates the results into a unified final proposal.

For example, a user request such as "Design an AI fitness app" would trigger the system to automatically generate a product requirements document, a technical architecture plan, a marketing and monetization strategy, a financial model with ROI projections, and a quality assurance review. The final output presents an integrated, multi-department project proposal that reflects the combined expertise of all participating agents.

## Agent Overview

| Agent     | Role               | Core Functions                                                                                                                    |
|-----------|--------------------|-----------------------------------------------------------------------------------------------------------------------------------|
| CEO Agent | Company Management | Interprets user requirements, defines strategic objectives, allocates tasks, coordinates agent workflow, integrates final outputs |

| Agent             | Role                     | Core Functions                                                                                                                                 |
|-------------------|--------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|
| Product Agent     | Product Manager          | Defines product features, analyses user needs, conducts competitive analysis, produces Product Requirements Documents (PRD)                    |
| Engineering Agent | Engineering Team         | Designs system architecture, evaluates technical feasibility, defines data flows, produces implementation roadmap                              |
| QA Agent          | Quality Assurance Team   | Reviews agent outputs for clarity, consistency, and completeness; performs cross-agent checking; produces structured QA review report          |
| Finance Agent     | Finance & Marketing Team | Analyses target market, formulates go-to-market strategy, evaluates costs, models ROI and break-even, produces commercial and financial report |

### Expected Outcomes

The expected outcome of this system is a fully automated, multi-perspective business proposal generation platform. By simulating cross-departmental collaboration through AI agents, the system reduces the time and effort required to produce comprehensive project plans. The system also improves the quality and consistency of generated proposals by incorporating structured quality assurance review and centralized orchestration.

# Scope of Work

---

## CEO Agent

### Role and Responsibilities

The CEO Agent serves as the central orchestration and decision-making component of the system. It interprets user requirements, decomposes complex objectives into structured subtasks, allocates tasks to specialized agents, coordinates cross-agent workflows, and integrates all outputs into a coherent final proposal. Its role is analogous to executive leadership in a real-world organization.

### Core Functions

- **Requirement Interpretation:** Analyses unstructured user input and transforms it into a structured project definition, identifying the functional domains required and expected deliverables.
- **Task Decomposition and Allocation:** Decomposes the overall objective into subtasks aligned with departmental responsibilities and assigns each to the appropriate specialized agent.
- **Workflow Orchestration:** Manages the sequence of agent interactions and handles inter-agent dependencies, ensuring upstream outputs are available to downstream agents when required.
- **Result Aggregation:** Consolidates and synthesizes all agent outputs into a unified project proposal, resolving inconsistencies across departmental perspectives.
- **Governance and Explainability:** Logs all task assignments and decision steps to maintain traceable reasoning paths and support transparency and accountability.
- **Security:** Enforces role-based information sharing between agents and validates agent outputs before integration into the final proposal.

### LLMOps and MLOps

The CEO Agent operates within an LLMOps-enabled infrastructure that supports logging of agent interactions, monitoring of system performance, and version control of prompts and agent configurations. These practices ensure reproducibility, observability, and continuous improvement of multi-agent collaboration.

## Product Agent

### Role and Responsibilities

The Product Agent acts as the Product Manager within the system. It ingests the CEO Agent's strategic direction, analyses user needs, conducts competitive analysis, and produces structured Product Requirements Documents (PRDs) for consumption by the Engineering Agent and other downstream agents.

## Core Functions

- Requirement Analysis: Parses JSON-formatted strategic mandates from the CEO Agent and translates them into product specifications.
- Competitive Analysis: Autonomously triggers web search tools to gather market intelligence and inform feature prioritization.
- PRD Generation: Produces markdown-formatted PRDs with standardized acceptance criteria, including Chain-of-Thought (CoT) logging to provide transparent rationale for each feature decision.
- Human-in-the-Loop (HITL) Checkpoint: A human stakeholder reviews and approves the finalized PRD before handoff to the Engineering Agent, with fallback to predefined templates if the LLM fails to generate a viable document.

## Assurance, Trust, Fairness, and Ethics

The Product Agent incorporates guardrails to ensure inclusive design, avoiding discriminatory assumptions regarding gender, race, or accessibility in user personas and feature definitions. PII provided in user input is masked via data sanitization pipelines before processing to prevent prompt injection and data leakage. Role-Based Access Control (RBAC) is enforced for prompt modification.

## LLMOps and MLOps

All prompt iterations are version-controlled. Telemetry integration (e.g., LangSmith) monitors token consumption, latency, and API performance. Evaluation metrics such as RAGAS assess the factual accuracy and comprehensiveness of generated PRDs against the original user prompt.

## Engineering Agent

### Role and Responsibilities

The Engineering Agent serves as the technical expert in the system. It analyses product requirements from the Product Agent and designs a technically feasible system architecture, including backend services, AI model integration, databases, and APIs.

## Core Functions

- Technical Requirement Analysis: Translates product requirements into technical specifications, identifying infrastructure needs and integration requirements.
- System Architecture Design: Designs modular backend architecture with clear API specifications, scalability planning, and risk identification.
- Technical Documentation: Produces structured documentation with traceable links between product requirements and engineering decisions.
- Cross-Agent Input: Provides cost and feasibility inputs to the Finance Agent and technical constraints to the Product Agent.

## Assurance, Trust, Fairness, and Ethics

The Engineering Agent incorporates responsible AI principles into system design, including secure data handling, bias awareness in AI models, privacy protection through encryption and anonymisation, and compliance with ethical AI guidelines. Architecture decisions are accompanied by structured explanations to improve stakeholder trust and explainability.

## Security and Governance

Security measures include authentication and authorisation mechanisms, secure communication protocols between agents, encryption of sensitive data, and role-based access control. Governance is supported through version control, architecture documentation, and system logs that enable auditing and future maintenance.

## LLMOps and MLOps

The Engineering Agent integrates AI components with defined MLOps/LLMOps pipelines for model invocation, prompt management, monitoring, and logging of agent interactions to ensure reproducible and observable system behaviour.

## QA Agent

### Role and Responsibilities

The QA Agent reviews the quality of outputs produced by the Product Agent, Engineering Agent, and Finance Agent before final integration by the CEO Agent. It does not generate business content; instead, it evaluates whether existing outputs are clear, complete, consistent, and aligned with the user's original request.

### Core Functions

- Individual Output Review: Assesses each agent's output for clarity, relevance, and internal coherence against predefined review criteria (clarity, consistency, completeness, feasibility, and alignment).
- Cross-Agent Consistency Checking: Identifies contradictions, missing connections, unsupported claims, and gaps between outputs produced by different agents.
- Risk Identification: Flags unclear requirements, unrealistic plans, weak justifications, or implementation gaps that may affect the quality of the final proposal.
- QA Review Report: Produces a structured report with prioritised issues and improvement recommendations for the CEO Agent.

### Accountability and Governance

The QA Agent applies consistent review criteria across all outputs to ensure fair and reliable feedback. Final decisions on acceptance, revision, or integration remain with the CEO Agent. The QA Agent provides issue explanations to support transparent, actionable feedback, and avoids overconfident judgement when available information is incomplete.

## Finance Agent

## **Role and Responsibilities**

The Finance Agent serves as the combined Finance and Marketing department within the system. It analyses the target market, devises go-to-market (GTM) strategies, evaluates costs derived from product and engineering plans, and produces financial models including ROI projections and break-even analysis. This dual scope ensures the final proposal is both commercially viable and financially sustainable.

## **Core Functions**

- **Market Strategy:** Defines target user groups, analyses market positioning, identifies promotional channels, and determines the product's commercial value proposition.
- **Financial Modelling:** Evaluates infrastructure, API, and marketing costs (including Customer Acquisition Cost); models revenue projections based on anticipated market capture and pricing strategy; produces break-even analysis.
- **Commercial Report:** Produces a comprehensive report covering GTM strategy, budget breakdowns, cash flow forecasts, and commercial risks (e.g., market sizing constraints relative to engineering costs).

## **Internal Consistency and Explainability**

The Finance Agent ensures that marketing strategies are financially consistent with overall budget constraints. It transparently documents the reasoning behind channel selection, pricing assumptions, and financial formulas used to calculate ROI and break-even points, ensuring the business logic is auditable by the CEO Agent.

## Effort Estimates

---

The following Work Breakdown Structure (WBS) summarises the estimated development effort for the AI CompanyOps Multi-Agent System. Tasks are aggregated across all agents and system-level components to reflect the full implementation scope.

### CEO Agent

| Task                                            | Description                                                                                                                                                                        | Estimated Effort |
|-------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|
| Requirement Interpretation & Strategic Planning | Design and implement reasoning logic to interpret user input and transform high-level requests into structured project objectives, including prompt templates and goal extraction. | 3 days           |
| Task Decomposition & Agent Assignment           | Develop mechanisms to decompose project objectives into structured subtasks aligned with departmental responsibilities and define standardized task formats for agent handoff.     | 2 days           |
| Agent Orchestration & Workflow Coordination     | Implement orchestration logic managing agent interaction order, dependency handling, and information flow between upstream and downstream agents.                                  | 4 days           |
| Result Aggregation & Strategic Integration      | Develop consolidation mechanisms to synthesize multi-agent outputs into a unified proposal, resolving cross-agent inconsistencies.                                                 | 3 days           |
| Governance, Explainability & Monitoring         | Implement logging of agent interactions, traceable reasoning paths, and LLMOps monitoring for performance observability.                                                           | 2 days           |
| Security & Integration Validation               | Implement role-based information sharing, output validation, and end-to-end integration testing across all agents.                                                                 | 2 days           |

### Product Agent

| Task                               | Description                                                                                                                                                                                     | Estimated Effort |
|------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|
| Requirements & Architecture Design | Define Product Agent scope, input/output schemas (JSON/Markdown), AI governance guardrails (PII masking, bias mitigation), and HITL approval workflows.                                         | 4 days           |
| Core LLM & Prompt Engineering      | Draft, test, and refine the core system prompt; implement Chain-of-Thought logging; integrate external tools (e.g., Web Search API for competitive analysis); establish prompt version control. | 5 days           |



| Task                               | Description                                                                                                                                                                      | Estimated Effort |
|------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|
| System Integration & Orchestration | Build data pipelines: CEO Agent → Product Agent (input parsing) and Product Agent → Engineering Agent (PRD handoff); automate task creation.                                     | 3 days           |
| LLMOps, Testing & Deployment       | Set up telemetry monitoring (token tracking, cost, latency); implement schema validation tests; conduct red teaming and prompt injection testing; produce security audit report. | 3 days           |

## Engineering Agent

| Task                                                    | Description                                                                                                                                                                       | Estimated Effort |
|---------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|
| Technical Requirement Analysis & Feasibility Study      | Analyse functional requirements from other agents and translate them into technical specifications, identifying system constraints and infrastructure requirements.               | 3 days           |
| System Architecture & Agent Communication Design        | Design overall system architecture including backend services, database structure, and AI model integration; define communication protocols and structured data exchange formats. | 4 days           |
| Engineering Agent Development & AI Pipeline Integration | Implement technical reasoning logic and workflow; integrate AI components; define MLOps/LLMOps pipelines for model invocation, prompt management, and interaction logging.        | 5 days           |
| System Integration, Security & Documentation            | Conduct integration testing with other agents; implement authentication and access control; produce technical documentation of system architecture and implementation.            | 3 days           |

## QA Agent

| Task                                | Description                                                                                                                                                                                           | Estimated Effort |
|-------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|
| Review Criteria & Logic Design      | Define evaluation criteria (clarity, consistency, completeness, feasibility, alignment) and design the review logic for both individual agent output review and cross-agent consistency checking.     | 5 days           |
| Agent Implementation & Integration  | Implement the QA Agent's review workflow; integrate into the multi-agent pipeline so it receives outputs from Product, Engineering, and Finance Agents and passes the review report to the CEO Agent. | 5 days           |
| Testing, Refinement & Documentation | Test QA Agent performance against sample multi-agent outputs; refine review criteria and output formatting; produce technical documentation.                                                          | 4 days           |

## Finance Agent

| Task                                            | Description                                                                                                                                                                                                                   | Estimated Effort |
|-------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|
| Business Logic & GTM Framework Design           | Define the combined marketing and financial analysis scope; design GTM frameworks, cost models, and ROI calculation methodology; establish internal consistency rules between marketing strategy and financial constraints.   | 5 days           |
| Agent Implementation & Prompt Engineering       | Implement the Finance Agent's dual-domain reasoning workflow; engineer prompts that balance creative marketing generation with rigorous financial computation; integrate input pipelines from Product and Engineering Agents. | 6 days           |
| Hallucination Mitigation, Testing & Integration | Validate that the LLM does not hallucinate financial calculations; conduct cross-agent integration testing; verify alignment of commercial report outputs with CEO Agent consolidation requirements.                          | 4 days           |